

A Compositional Deadlock Detector for Android Java

Darius Mureşan Henrique Luz Rui Xavier

November 11, 2025

Aarhus University
Department of Computer Science

Overview

Overview

- We present a **compositional static analysis** for detecting deadlocks in Android Java.
- Designed for **industrial-scale** codebases (tens of millions of lines).
- Implemented in **INFER**, deployed at Facebook for 2+ years.
- Achieves a **54% developer fix rate**.

Motivation

Why Study Deadlocks?

- Deadlocks are a key challenge in concurrent programming.
- Occur when threads **cyclically wait on each other's locks**.
- Result: entire system halts — loss of responsiveness, reliability.
- Especially relevant for **Android apps** using Java's **synchronized blocks**.

Classical Example

Dijkstra's Dining Philosophers

Five philosophers share forks and a bowl of spaghetti. Each must hold two forks to eat — potential for a circular wait.

- Illustrates the essential nature of deadlock: mutual waiting.
- Analogous to threads waiting for locks.

Deadlocks in Industry

- At Facebook, Android apps exceed **10M+ lines of code**.
- Thousands of commits daily ⇒ rapid iteration.
- Developers need feedback in **under 15 minutes**.
- Whole-program analyses are too slow and memory-intensive.

Goal

Fast, scalable, accurate deadlock detection integrated into CI pipelines.

Challenges in Traditional Analyses

- Reanalyze entire program on each commit.
- Poor scalability and high false-positive rates.
- Lack **compositionality**.
- Need: an analysis that focuses only on changed code and its dependencies.

Research Gap and Question

The Research Gap

- Existing tools are **non-compositional**.
- They often trade soundness for scalability.
- We need a **mathematically grounded, incremental** analysis.

Core Challenge

Can we design a deadlock detector that is both **sound in theory** and **scalable in practice**?

Research Question

Main Question

Can we develop a **compositional deadlock detector** for **Android Java** that is sound, complete, and efficient for large industrial codebases?

- Model Java concurrency with **balanced, re-entrant locks**.
- Detect deadlocks via **critical pairs**.
- Integrate into **INFER** for continuous integration.

Concurrent Programs

Abstract Language

- Simplified model of Java concurrency:

$$C ::= \text{skip} \mid p() \mid \text{acq}(\ell) \mid \text{rel}(\ell) \mid C; C \mid \text{if}(*) \text{ then } C \text{ else } C \mid$$

- Non-recursive procedures.
- **Balanced statements:** every acq has matching rel.
- Parallel program: $C_1 \parallel C_2 \parallel \dots \parallel C_n$

Balanced (Nested) Locks

- Threads acquire and release locks in **LIFO** order.
- Corresponds to `synchronized` in Java.
- Ensures locks are **re-entrant and scoped**.

Example

acq(x); acq(y); rel(y); rel(x) **Balanced**
acq(x); rel(y); rel(x) **Unbalanced**

- Each thread has a **lock state**: mapping locks → acquisition counts.
- Deadlock occurs when every thread can take a local step, but no joint step is possible.

$$\langle C_1 || C_2, (L_1, L_2) \rangle$$

Key Idea

Deadlock arises when threads hold disjoint locks yet each waits on a lock held by another.

Program Execution Traces

Executions as Traces

- Executions can be represented as strings of lock acquisitions/releases — e.g. ‘x y y x’.
- Balanced programs produce **Dyck words** (well-nested parentheses).
- Captures the essential locking behaviour.

$$L(C) = \{\text{all possible lock traces of } C\}$$

Example Trace

Example

```
acq(x); if(*) then acq(y); rel(y); else acq(z);  
rel(z); rel(x)
```

$$L(C) = \{ xyyx, xzzx \}$$

- Balanced structure guarantees decidability.
- Enables abstract reasoning about possible interleavings.

Traces as Finite Automata

- Each balanced statement can be viewed as a **finite automaton** over lock actions.
- Allows algorithmic computation of lock dependencies.
- Provides the foundation for critical-pair analysis.

Soundness and Completeness

Critical Pair

(X, ℓ) : some execution acquires lock ℓ while holding all locks in X .

- Captures possible lock dependencies.
- Computed for each sequential thread.

Key Theorem (Simplified)

Theorem 4.4 — Deadlock Condition

Program $C_1||\dots||C_n$ deadlocks iff there exist critical pairs (X_i, ℓ_i)
s.t.

$$\ell_i \in \bigcup_{j \neq i} X_j \quad \text{and} \quad X_i \cap \bigcup_{j \neq i} X_j = \emptyset$$

Intuition: Each thread holds a lock another needs.

Illustrative Example

Two-Thread Example

C1: acq(x); acq(y); rel(y); rel(x) C2: acq(y);
acq(x); rel(x); rel(y)

$$\text{Crit}(\text{C1}) = \{(\emptyset, x), (\{x\}, y)\}$$

$$\text{Crit}(\text{C2}) = \{(\emptyset, y), (\{y\}, x)\}$$

Since $x \in \{y\}$ and $y \in \{x\}$, the condition holds — **deadlock!**

Proof Structure

- Define executions \rightarrow and parallel composition.
- Show equivalence between execution semantics and trace semantics.
- Prove critical-pair condition is **sound** (no missed deadlocks) and **complete** (no spurious ones).

Result

Existence of deadlock $\models_{\mathcal{L}}$ conflict between critical pairs.

Complexity

Computing Critical Pairs

- Recursive equations (C1–C6) compute $\text{Crit}(C)$ compositionally.
- Each construct (if, while, seq) has a local combination rule.
- **Example:**

$$\text{Crit}(\text{acq}(\ell); C; \text{rel}(\ell)) = \{(\emptyset, \ell)\} \cup \{(X \cup \{\ell\}, \ell') \mid (X, \ell') \in \text{Crit}(C)\}$$

Complexity Results

- **Finite and computable:** $\text{Crit}(C)$ always finite.
- Deadlock detection problem is **decidable** and in NP.
- Non-recursive programs \Rightarrow quadratic time.
- With procedures \Rightarrow quasi-exponential.

Implementation

Implementation Overview

- Implemented as an **abstract interpretation** within INFER.
- Computes method summaries: critical pairs + thread identity.
- Compositionally reuses summaries of unchanged methods.

Core Idea

Analyse only modified methods and their dependents.

Abstract State Representation

$\alpha = \langle L, Z \rangle$ where L = lock state, Z = set of critical pairs

- Join operation: $\langle L, Z_1 \rangle \sqcup \langle L, Z_2 \rangle = \langle L, Z_1 \cup Z_2 \rangle$
- Each command updates this abstract state.

Compositionality

- Procedure call depends only on:
 1. Current abstract state.
 2. Precomputed summary of the callee.
- Enables incremental reanalysis — ideal for CI/CD.

$$Jp()K\langle L, Z \rangle = \langle L, Z \cup f(L, \text{Crit}(\text{body}(p))) \rangle$$

- **Balanced locking:** uses synchronized.
- **Partial path sensitivity:** e.g., for tryLock() and UI threads.
- **Lock naming:** access-path abstraction (this.f.g, etc.).
- **Thread inference:** uses annotations like @UiThread.

Deployment and Results

Industrial Deployment

- Deployed as part of Facebook's **continuous integration system**.
- Runs automatically on every Android commit.
- Appears as an automated "reviewer" commenting on potential deadlocks.

Results and Impact

- Deployed for **2+ years** on all Android commits.
- **500+** deadlock reports issued.
- **54%** of reports fixed by developers.
- Median runtime (all analyses): **90 s per commit.**
- Analyses **2k–5k methods per commit.**

Conclusion and Related Work

Related Work

- Builds on automata-theoretic analyses of **pushdown systems**.
- Compared to prior tools:
 - Compositional, not whole-program.
 - Targets balanced re-entrant locks.
 - Prioritizes **actionable results** over completeness.

Conclusion

- Developed a **sound and complete** compositional analysis.
- Scales to tens of millions of lines.
- Successfully deployed in industry with tangible impact.
- Formalized and proven in **Coq (8.7k LOC)**.

Future Work

Extend to recursive calls, deterministic control, and nested parallelism.

Thank you!

Questions?