

A Compositional Deadlock Detector for Android Java

Darius Muresan

Henrique Luz

Rui Xavier

Aarhus University
Department of Computer Science

November 11, 2025



AARHUS UNIVERSITY

Why Study Deadlocks?

- Deadlocks are a fundamental challenge in concurrent programming.
- They occur when threads cyclically wait on each other's locks, preventing further progress.
- Detecting and preventing them is critical for reliability, especially in large-scale systems.
- In industrial settings (e.g., Android applications at Facebook), codebases can exceed tens of millions of LoC, making traditional whole-program analyses infeasible.

The Industrial Challenge

- Modern software development involves:
 - Massive, continuously evolving codebases.
 - Frequent commits and rapid code reviews.
 - Strong requirements for developer feedback in under 15 minutes.
- Conventional static analyses:
 - Require analyzing entire programs.
 - Are too slow or memory-intensive for this context.
 - Often produce many false positives.

Goal

Enable fast, scalable, and accurate deadlock detection that integrates seamlessly into continuous integration pipelines.

The Research Gap

- Most existing tools:
 - Lack compositionality by reanalyzing the whole program each time.
 - Sacrifice scalability for soundness.
- Desired: a compositional, mathematically grounded approach that can analyze only changed code and its dependents.

Core Challenge

How can we design a deadlock analysis that is both **theoretically sound** and **practically scalable** to millions of lines of concurrent Android Java code?

Our Research Question

Research Question

Can we develop a **compositional deadlock detector** for **Android Java** that is sound, complete, and efficient enough to run on commercial-scale codebases under active development?

- Key aspects explored:
 - Abstract modeling of real Java concurrency (balanced, re-entrant locks).
 - Deadlock characterization via **critical pairs**.
 - Compositional analysis integrated into the **INFER** framework.
- Goal: bridge the gap between **theory** and **industrial-scale deployment**.

Motivation
oooo

Concurrent Programs
●ooo

Program Executions
ooooo

Soundness and Completeness
oooooo

Complexity
oo

Implementation
○

Conclusion and Related Work
ooo

Concurrent Programs

How to look at Concurrent Programs?

```
type α node = {  
    data : α;  
    mutable prev : α node;  
    mutable next : α node;  
}
```

Gospel + OCaml

```
Definition Node A (v: A) (n p c: loc) : hprop :=  
c ↤ '{ data' := v; next' := n; prev' := p }.
```

CFML

Inner Node

```
type α innerNode =  
| Nil  
| Cons of α node
```

Gospel + OCaml

```
Definition InnerNode A (L: list A) (p: innerNode_ A) : hprop :=  
match L with  
| [] => [p = Nil]  
| _ => ∃ (c q: loc), [p = Cons c] * c ↣ NodeSeg L c q q  
end.
```

CFML

Doubly Linked List

```
type α dblist = {  
    mutable head : α innerNode;  
    mutable tail : α innerNode;  
    mutable length : int;  
}
```

Gospel + OCaml

Definition Dblist A (L: list A) (l: loc) : hprop :=
 $\exists (p q: \text{innerNode_ } A),$
 $(l \mapsto \{\text{head' := } p; \text{tail' := } q; \text{length' := } \text{length } L\}) \star$
 $(\text{If } L = [] \text{ then } [p = \text{Nil}] \star [q = \text{Nil}]$
 else
 $\exists x L' h t, [L = x :: L'] \star [p = \text{Cons } h] \star [q = \text{Cons } t]$
 $\star (h \rightsquigarrow \text{NodeSeg } L h t t)).$

CFML

Program Executions and their traces

```
val create : unit → α dblist
val is_empty : α dblist → bool
val clear : α dblist → unit
val remove_head : α dblist → α node
val remove_tail : α dblist → α node
val reverse : α dblist → unit
val append : α dblist → α dblist → α dblist
val josephus : α dblist → int → unit
val fold_right : (α → β → β) → α dblist → β → β
val fold_left : (α → β → α) → α → β dblist → α
val iter_right : (α → β) → α dblist → β
val iter_left : (α → β) → α dblist → unit
...

```

GOSPEL + OCaml

```
val create : unit → α dblist
val is_empty : α dblist → bool
val clear : α dblist → unit
val remove_head : α dblist → α node
val remove_tail : α dblist → α node
val reverse : α dblist → unit
val append : α dblist → α dblist → α dblist
val josephus : α dblist → int → unit
val fold_right : (α → β → β) → α dblist → β → β
val fold_left : (α → β → α) → α → β dblist → α
val iter_right : (α → β) → α dblist → β
val iter_left : ('a -> 'b) -> 'a dblist -> unit
...

```

GOSPEL + OCaml

Predicates

Filliâtre and Pereira [6] introduced two key logical predicates, *permitted* and *complete*, which allow us to reason both about the correctness of an iterator's behavior and its termination.

Permitted – Visited Elements

$$\text{permitted}(v, s) \triangleq ||v|| \leq ||s|| \wedge \forall i : 0 \leq i \leq ||v|| \implies v[i] = s[i]$$

Complete – Iteration Termination

$$\text{complete}(v, s) \triangleq ||v|| = ||s||$$

iter_left Function

```
val iter_left : ( $\alpha \rightarrow \beta$ )  $\rightarrow \alpha$  dblist  $\rightarrow$  unit Gospel + OCaml
(*@ r = iter_left f collection
  iterspec
  ~permitted: (fun v  $\rightarrow$  length v  $\leq$  length collection  $\wedge$ 
               $\forall i. 0 \leq i < \text{length } v \rightarrow v[i] = \text{collection}[i]$ )
  ~complete: (fun v  $\rightarrow$  length v = length collection) *)
```

This specification is a recent addition to Gospel by Ion Chirica and Mário Pereira [3].

Soundness and Completeness

OCaml Code

```
let rec iter_left_aux f node tail =
  f node.data;
  if not (tail == node) then
    iter_left_aux f node.next tail

let iter_left f db =
  match db.head with
  | Nil → ()
  | Cons n →
    match db.tail with
    | Nil → assert false
    | Cons tail → iter_left_aux f n tail
```

OCaml

Consumer Function Specification

$$\forall x L1 L2, L = L1 ++ x :: L2 \rightarrow$$

SPEC (f x)
PRE (I L1)
POSTUNIT (I (L1 & x))

CFML

iter_left Specification

```
Lemma Triple_iter_left : ∀ (I: list A → hprop) L (f: val) t,  
  (∀ x L1 L2, L = L1 ++ x :: L2 →  
   SPEC (f x)  
   PRE (I L1)  
   POSTUNIT (I (L1 & x))) →  
  SPEC (iter_left f t)  
  PRE (t ↪ Dblist L ∗ I [])  
  POSTUNIT (t ↪ Dblist L ∗ I L).
```

CFML

iter_left_aux Specification


To verify the higher-level `iter_left` function, we require a more fine-grained specification for its auxiliary, lower-level implementation.

Lemma `Triple_iter_left_aux` :

CFML

```
  ∀ (f: val) (I : list A → hprop) (L L1 L2: list A) n e h p,  
    (∀ x L1 L2, L = L1 ++ x :: L2 →  
      SPEC (f x)  
      PRE (I L1)  
      POSTUNIT (I (L1 & x))) →  
    L = L1 ++ L2 →  
    L2 ≠ nil →  
    SPEC (iter_left_aux f n e)  
    PRE (n ↦ NodeSeg L2 h e p * I L1)  
    POSTUNIT (n ↦ NodeSeg L2 h e p * I L).
```

Upper Complexity Bounds



Implementation

Motivation
oooo

Concurrent Programs
oooo

Program Executions
oooooo

Soundness and Completeness
oooooo

Complexity
oo

Implementation
o

Conclusion and Related Work
●oo

Conclusion and Related Work

Conclusions

- Designed and implemented a circular doubly linked list in OCaml.
- Verified key operations (such as `create`, `clear`, `iter_left`, `fold_left`) using CFML.
- Demonstrated feasibility of verifying mutable data structures in a functional language.
- Identified limitations within the current toolchain.

Reflections & Future Work

Reflections

- Formal verification remains time-intensive (some proofs may take months).
- Highlights both promise and current challenges in formally verifying real-world functional code.

Future Work

- Verify the remaining operations and auxiliary lemmas.
- Perform the refinement by comparing manual specs with Peter-generated ones.

Bibliography I

- [1] Arthur Charguéraud. “Characteristic formulae for the verification of imperative programs”. In: *SIGPLAN Not.* 46.9 (Sept. 2011), pp. 418–430. ISSN: 0362-1340. DOI: 10.1145/2034574.2034828. URL:
<https://doi.org/10.1145/2034574.2034828>.
- [2] Arthur Charguéraud et al. “GOSPEL - Providing OCaml with a Formal Specification Language”. In: *Formal Methods - The Next 30 Years - Third World Congress, FM 2019, Porto, Portugal, October 7-11, 2019, Proceedings*. Ed. by Maurice H. ter Beek, Annabelle McIver, and José N. Oliveira. Vol. 11800. Lecture Notes in Computer Science. Springer, 2019, pp. 484–501. DOI: 10.1007/978-3-030-30942-8__29. URL:
https://doi.org/10.1007/978-3-030-30942-8%5C_29.

Bibliography II

- [3] Ion Chirica and Mário Pereira. “Unfolding Iterators: Specification and Verification of Higher-Order Iterators, in OCaml”. In: *arXiv preprint arXiv:2506.20310* (2024). DOI: 10.48550/arXiv.2506.20310. arXiv: 2506.20310 [cs.PL]. URL: <https://arxiv.org/abs/2506.20310>.
- [4] Michael R. Clarkson et al. *OCaml Programming: Correct + Efficient + Beautiful*. <https://cs3110.github.io/textbook/cover.html>. 2025.
- [5] Jean-Christophe Filiâtre and Clément Pascutto. “Ortac: Runtime Assertion Checking for OCaml (Tool Paper)”. In: *Runtime Verification: 21st International Conference, RV 2021, Virtual Event, October 11–14, 2021, Proceedings*. Berlin, Heidelberg: Springer-Verlag, 2021, pp. 244–253. ISBN: 978-3-030-88493-2. DOI: 10.1007/978-3-030-88494-9_13. URL: https://doi.org/10.1007/978-3-030-88494-9_13.

Bibliography III

- [6] Jean-Christophe Fillâtre and Mário Pereira. “A Modular Way to Reason About Iteration”. In: *NASA Formal Methods - 8th International Symposium, NFM 2016, Minneapolis, MN, USA, June 7-9, 2016, Proceedings*. Ed. by Sanjai Rayadurgam and Oksana Tkachuk. Vol. 9690. Lecture Notes in Computer Science. Springer, 2016, pp. 322–336. DOI: 10.1007/978-3-319-40648-0_24. URL: https://doi.org/10.1007/978-3-319-40648-0%5C_24.
- [7] Marc Hermes and Robbert Krebbers. “Modular Verification of Intrusive List and Tree Data Structures in Separation Logic”. In: *15th International Conference on Interactive Theorem Proving (ITP 2024)*. Ed. by Yves Bertot, Temur Kutsia, and Michael Norrish. Vol. 309. Leibniz International Proceedings in Informatics (LIPIcs). Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2024, 19:1–19:18. ISBN: 978-3-95977-337-9. DOI: 10.4230/LIPIcs.ITP.2024.19. URL: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.ITP.2024.19>.

Bibliography IV

- [8] Mário Pereira and António Ravara. “Cameleer: A Deductive Verification Tool for OCaml”. In: July 2021, pp. 677–689. ISBN: 978-3-030-81687-2. DOI: 10.1007/978-3-030-81688-9_31.
- [9] Mário Pereira Rui Xavier and Tiago Soares. *Formal Verification of OCaml Programs with Dynamic Memory*. Companion Artifact. 2025. URL: <https://github.com/RuiXavier/adc-artifact-companion> (visited on 06/23/2025).
- [10] Tiago Lopes Soares, Ion Chirica, and M'ario Pereira. “Static and Dynamic Verification of OCaml Programs: The Gospel Ecosystem”. In: *Leveraging Applications of Formal Methods, Verification and Validation. Specification and Verification*. Ed. by Tiziana Margaria and Bernhard Steffen. Cham: Springer Nature Switzerland, 2025, pp. 247–265. ISBN: 978-3-031-75380-0.