

A Compositional Deadlock Detector for Android Java

Darius Mureşan Henrique Luz Rui Xavier

November 19, 2025

Aarhus University
Department of Computer Science

Motivation

Deadlocks: A Refresher

- A **deadlock** occurs when two or more threads are each waiting for the other to release a lock, so none can make progress.

Example (two threads, two locks)

Thread 1: acq(x); acq(y);

Thread 2: acq(y); acq(x);

- One possible execution:
 1. Thread 1 acquires x
 2. Thread 2 acquires y
 3. Thread 1 now waits for y, Thread 2 waits for x
- Neither thread can continue \implies **deadlock**.

Why Deadlocks Matter

- Concurrency is everywhere in Android apps:
 - UI thread + many background threads
 - Callback-driven, event-based execution
 - Shared mutable state (objects used as locks)
- A **deadlock** freezes part (or all) of the app:
 - UI becomes unresponsive
 - Background tasks never finish
 - Users perceive the app as *broken*
- Avoiding and detecting deadlocks is a core reliability problem.

Industrial Setting: Facebook Android Apps

- Target: large Android apps under continuous development:
 - Tens of millions of lines of code (LoC)
 - Thousands of revisions per day
- Goal of the analysis:
 - Run at code-review time on every commit
 - Give feedback to developers **within minutes**
- This rules out whole-program, from-scratch analyses for every change.

Key Challenges

- **Scalability:**
 - Cannot re-analyse the whole app for each commit
 - Need a compositional, incremental analysis
- **Usefulness for developers:**
 - Too many false positives ⇒ warnings are ignored
 - The focus is on **actionable reports**, not proving absence of deadlocks
- **The research question:** can we get both **theoretical guarantees** and **practical performance** at this scale?

Research Gap and Question

Limitations of Existing Approaches

- Many existing deadlock analyses:
 - Assume access to the **whole program**
 - Do not support compositional, change-focused analysis
- Tools often prioritise:
 - Strong soundness guarantees
 - But with high false-positive rates in practice
- For huge codebases, this is misaligned with developer needs:
 - Developers prefer fewer, more trustworthy warnings
 - It is acceptable to miss some rare deadlocks

Research Question

Main Question

Can we design a **compositional** static analysis for **Android Java** that detects deadlocks in large codebases, with:

- A clean **theoretical characterisation** of deadlocks
- A **decidable** and **tractable** core problem
- A practical implementation integrated in CI

High-Level Contributions

High-Level Contributions

- Abstract language with balanced re-entrant locks and nondeterministic control
- New deadlock condition based on **critical pairs** of threads
- Proof that deadlock detection in this language is decidable and in NP
- Compositional implementation in Facebook's INFER, deployed at scale

Critical Pairs and Deadlock Detection

Critical Pair (Informal)

A **critical pair** of a thread is a pair (X, ℓ) such that some execution of the thread acquires an *unheld* lock ℓ while already holding exactly the set of locks X .

- Captures **which locks are held** when a new lock is acquired.
- Abstracts away the concrete control-flow and interleavings.
- Each thread has a **finite** set of critical pairs.

Critical Pairs: A Simple Example

Statement

$$C = \text{acq}(x); \text{ acq}(y); \text{ rel}(y); \text{ rel}(x)$$

Computing $\text{Crit}(C)$

$$\text{Crit}(C) = \{ (\emptyset, x), (\{x\}, y) \}$$

- At the first acquisition of x : no locks held $\Rightarrow (\emptyset, x)$
- Later, the thread holds $\{x\}$ when acquiring $y \Rightarrow (\{x\}, y)$

Two-Thread Deadlock Condition

Deadlock Condition (2 Threads)

For two statements C_1 and C_2 running concurrently:

$$C_1 \parallel C_2 \text{ deadlocks} \iff \exists (X_1, \ell_1) \in \text{Crit}(C_1), (X_2, \ell_2) \in \text{Crit}(C_2)$$

such that

$$\ell_1 \in X_2, \quad \ell_2 \in X_1, \quad X_1 \cap X_2 = \emptyset.$$

- Each thread holds a lock the other is trying to acquire.
- The currently held lock sets do not overlap.
- Generalises to arbitrarily many threads (Theorem 4.4 in the paper).

Example: Classic Two-Thread Deadlock

Two Threads

$C_1 : acq(x); acq(y); skip; rel(y); rel(x)$

$C_2 : acq(y); acq(x); skip; rel(x); rel(y)$

- Critical pairs:

$$\text{Crit}(C_1) = \{(\emptyset, x), (\{x\}, y)\}$$

$$\text{Crit}(C_2) = \{(\emptyset, y), (\{y\}, x)\}$$

- Take $(X_1, \ell_1) = (\{x\}, y)$ and $(X_2, \ell_2) = (\{y\}, x)$:
 - $\ell_1 = y \in X_2 = \{y\}$
 - $\ell_2 = x \in X_1 = \{x\}$
 - $X_1 \cap X_2 = \emptyset$
- Condition holds \Rightarrow deadlock is possible.

Guard Locks: Breaking the Deadlock

Adding a Guard Lock

$$C'_1 = \text{acq}(z); C_1; \text{rel}(z)$$

$$C'_2 = \text{acq}(z); C_2; \text{rel}(z)$$

- Now:

$$\text{Crit}(C'_1) = \{(\emptyset, z), (\{z\}, x), (\{z, x\}, y)\}$$

$$\text{Crit}(C'_2) = \{(\emptyset, z), (\{z\}, y), (\{z, y\}, x)\}$$

- Any potentially conflicting critical pairs share lock z :

$$\{z, x\} \cap \{z, y\} = \{z\} \neq \emptyset$$

- Deadlock condition fails \Rightarrow no deadlock.
- Intuition: z acts as a *guard lock* protecting the region.

Concurrent Programs

Abstract Language for Concurrency

- Abstracts away data and heap; focuses only on locks and control.
- Statements C are built from:

$$\begin{aligned} C ::= & \text{skip} \mid p() \mid \text{acq}(\ell); C; \text{rel}(\ell) \mid C; C \\ & \mid \text{if}(*) \text{ then } C \text{ else } C \mid \text{while}(*) \text{ do } C \end{aligned}$$

- **Balanced statements** enforce lock discipline:
 - Locks must be released in LIFO order
 - Models scoped constructs like synchronized and `std::lock_guard`
- A parallel program is an n -tuple $C_1 \parallel \dots \parallel C_n$.

Example Program in the Abstract Language

Example

$$C = acq(x); (\text{if } (*) \text{ then } acq(y); rel(y) \text{ else skip}); rel(x)$$

- This program:
 - Always acquires and releases x
 - Nondeterministically either:
 - acquires and releases y , or
 - does nothing in the branch
- Shows how nested acquisitions map to lock scopes in Java.

Lock States and Configurations

- Locks are **re-entrant**:
 - A thread can acquire the same lock multiple times
 - Lock state $L : \text{Locks} \rightarrow \mathbb{N}$ counts acquisitions
- We write $[L] = \{\ell \mid L(\ell) > 0\}$ for held locks.
- A configuration is a pair $\langle C, L \rangle$:
 - C : current statement
 - L : lock state
- A concurrent configuration:

$$\langle C_1 || \dots || C_n, (L_1, \dots, L_n) \rangle$$

where each thread has its own lock state.

Deadlock Definition (1)

- Sequential step: $\langle C, L \rangle \rightarrow \langle C', L' \rangle$.

Example

$\langle \text{acq}(x); \text{skip}; \text{rel}(x), L \rangle \rightarrow \langle \text{skip}; \text{rel}(x), L[x^{++}] \rangle$

(acquiring lock x increases its count)

- Parallel step: advance one thread, provided it does not grab a lock already held by another thread.
- Intuitively, we are deadlocked when:
 - Each thread can individually make a step (sequential)
 - But no *parallel* step is possible any more

Deadlock Definition (2)

- Formal definition (simplified):
 - A concurrent configuration is deadlocked if at least two threads are stuck in this way.
 - A program deadlocks if some reachable configuration is deadlocked.

Program Execution Traces

Executions as Traces

- Each step either:
 - Leaves the lock state unchanged
 - Acquires a lock ℓ
 - Releases a lock ℓ
- We record only lock actions as a **trace** over alphabet Σ :

$$\Sigma = \{\ell \mid \ell \in \text{Locks}\} \cup \{\bar{\ell} \mid \ell \in \text{Locks}\}$$

Example

- Statement: `acq(x); if(*) then acq(y); rel(y) else acq(z); rel(z); rel(x)`
- Possible traces: $xy\bar{y}x$ and $xz\bar{z}x$

Executions as Traces (Examples)

Example Statement

$$C = \text{acq}(x); \text{acq}(y); \text{rel}(y); \text{rel}(x)$$

- Valid trace:

$x y \bar{y} \bar{x}$

Let us now consider:

$$C' = \text{acq}(x); \text{acq}(y); \text{rel}(x); \text{rel}(y)$$

- Invalid trace (bad nesting):

$x y \bar{x} \bar{y} \Rightarrow \text{lock } x \text{ released while } y \text{ still held}$

Dyck Words and Balanced Locking

- For balanced statements, traces are **Dyck words**:
 - Well-parenthesised strings of opens and closes
 - Locks always released in reverse order of acquisition
- Key property:
 - For any balanced statement C , all traces in $L(C)$ are Dyck words.
 - Executions never underflow the lock stack.
- This structure is crucial for the decidability and the critical-pair characterisation of deadlocks.

Dyck Words and Balanced Locking (Example)

- For balanced statements, traces form **Dyck words**:
 - Perfectly nested lock acquisitions / releases
 - The lock stack never underflows

Valid Dyck Word

$$x \, y \, \bar{y} \, z \, \bar{z} \, \bar{x}$$

- Corresponds to nested Acquire–Release behaviour:

$$\textit{acq}(x); \textit{acq}(y); \textit{rel}(y); \textit{acq}(z); \textit{rel}(z); \textit{rel}(x)$$

Invalid Word

$$x \, y \, \bar{x} \, \bar{y}$$

- Releases x before y : violates LIFO discipline.

Languages of Statements

- Each statement C defines a language $L(C) \subseteq \Sigma^*$:
 - All traces of possible executions of C
 - Defined inductively:

$$L(\text{skip}) = \{\varepsilon\}$$

$$L(\text{rel}(\ell)) = \{\bar{\ell}\}$$

$$L(p()) = L(\text{body}(p))$$

$$L(C_1; C_2) = L(C_1) \cdot L(C_2)$$

$$L(\text{acq}(\ell)) = \{\ell\}$$

$$L(\text{while}(\ast) \text{ do } C) = L(C)^*$$

$$L(\text{if}(\ast) \text{ then } C_1 \text{ else } C_2) = L(C_1) \cup L(C_2)$$

- For balanced C , $L(C)$ is regular and consists of Dyck words.

Soundness and Completeness

Cumulative Lock Effect of a Trace

- For a trace $u \in \Sigma^*$, $\langle u \rangle$ denotes its **cumulative lock effect**:
 - Start from the empty lock multiset
 - Read actions in u left to right
 - Each x increments the count of lock x
 - Each \bar{x} decrements the count of lock x
- The set of *currently held locks* after reading u is:

$$[\langle u \rangle] = \{ \ell \mid \langle u \rangle(\ell) > 0 \}.$$

Example

Trace prefix: $u = x y \bar{y}$

$$\langle \varepsilon \rangle = \emptyset$$

$$\langle x \rangle = \{x\}$$

$$\langle xy \rangle = \{x, y\}$$

$$\langle xy\bar{y} \rangle = \{x\}$$

So:

$$\lfloor \langle u \rangle \rfloor = \{x\}.$$

Critical Pairs: Formal Definition

Critical Pairs of a Statement

For a balanced statement C :

$$\text{Crit}(C) = \{(\lfloor \langle u \rangle \rfloor, \ell) \mid \exists v. \, ulv \in L(C) \text{ and } \ell \notin \lfloor \langle u \rangle \rfloor\}$$

where $\langle u \rangle$ is the cumulative lock effect of trace u .

- This is equivalent to: *there exists an execution that acquires ℓ while holding exactly X .*
- The paper shows the execution-based and language-based definitions coincide.

Critical Pairs: Formal Definition (Example)

Example Trace

$$u\ell v = x y z \bar{z} \bar{y} \bar{x}$$

- Take $u = x y$, $\ell = z$:

$$\langle u \rangle = \{x, y\}$$

- Since $z \notin \{x, y\}$, acquiring z yields:

$$(\{x, y\}, z) \in \text{Crit}(C)$$

- Meaning: During some execution, the thread acquires z while holding exactly $\{x, y\}$.

Deadlock Characterisation Theorem

Theorem 4.4 (Simplified)

A parallel program $C_1||\dots||C_n$ deadlocks iff there exists an index set $I \subseteq \{1, \dots, n\}$ with $|I| \geq 2$ and critical pairs $(X_i, \ell_i) \in \text{Crit}(C_i)$ for each $i \in I$ such that:

$$X_i \cap \bigcup_{j \neq i} X_j = \emptyset \quad \text{and} \quad \ell_i \in \bigcup_{j \neq i} X_j \quad \text{for all } i \in I.$$

- Each thread holds locks needed by the others.
- Held-lock sets are pairwise disjoint.
- This condition is both **sound** and **complete**.

Complexity

Computing Critical Pairs Compositionally

- The paper gives 6 equations describing $\text{Crit}(C)$ by syntax of C .
- Examples:

$$\text{Crit}(\text{skip}) = \emptyset$$

$$\text{Crit}(p()) = \text{Crit}(\text{body}(p))$$

$$\text{Crit}(C; C') = \text{Crit}(C) \cup \text{Crit}(C')$$

$$\text{Crit}(\text{if } (*) \text{ then } C \text{ else } C') = \text{Crit}(C) \cup \text{Crit}(C')$$

$$\text{Crit}(\text{acq}(\ell); C; \text{rel}(\ell)) =$$

$$\{(\emptyset, \ell)\} \cup \{(X \cup \{\ell\}, \ell') \mid (X, \ell') \in \text{Crit}(C), \ell' \neq \ell\}$$

- These identities allow a bottom-up computation of $\text{Crit}(C)$.

Complexity Bounds

- **Finite and computable:**
 - For any balanced C , $\text{Crit}(C)$ is finite.
- Complexity of deadlock detection:
 - The deadlock problem for this language is **decidable** and lies in **NP**.
 - Idea: nondeterministically guess a set of threads and critical pairs, then check the deadlock condition in polynomial time.
- For programs without procedure calls:
 - Computing $\text{Crit}(C)$ is polynomial in program size.
- Lower bounds (e.g. NP-completeness) are left as future work.

Implementation

Core Abstract Interpretation

- Implementation is an abstract interpretation that computes critical pairs.
- Abstract state:

$$\alpha = \langle L, Z \rangle$$

where

- L : abstract lock state
- Z : set of (approximate) critical pairs
- Each command C defines a transformer $\llbracket C \rrbracket(\alpha)$.
- The join operation on states:

$$\langle L, Z_1 \rangle \sqcup \langle L, Z_2 \rangle = \langle L, Z_1 \cup Z_2 \rangle$$

Compositionality of the Analysis

- Procedure calls are handled via **summaries**:
 - For each procedure p we precompute $\text{Crit}(\text{body}(p))$.
 - At a call $p()$, we combine the current state with p 's summary.
- Consequences:
 - When code changes, only affected procedures and their callers need re-analysis.
 - Most of the program can be reused from previous runs.
 - This is essential for deployment in continuous integration.

Adapting to Android Java

- The abstract language is mapped to real Java/Android code:
 - synchronized methods/blocks ⇒ balanced lock regions
 - Re-entrant monitors modelled as nested acquisitions
- Android-specific refinements:
 - Partial path-sensitivity for methods like tryLock() and UI-thread checks
 - Lock naming via access-paths (this.f.g, etc.)
 - Thread identity domain (@UiThread, @WorkerThread, background)
- Implemented as the starvation analyser inside INFER.

Deployment and Results

Integration in Facebook's CI

- INFER is part of Facebook's continuous integration:
 - Every Android commit triggers static analyses, including deadlock analysis.
 - The analyser appears as an automated reviewer on code reviews.
- The deadlock analysis targets **code changes**, not whole apps:
 - Summarise modified methods and their dependents
 - Use heuristics to find relevant methods that share locks

Quantitative Impact

- Deployed on all Android commits for ~2 years.
- Scale:
 - Hundreds of thousands of commits analysed
 - Typically 2k – 5k methods per commit
- Performance:
 - Median analysis time of 90 seconds per commit
 - Average analysis time of 213 seconds per commit
- Effectiveness:
 - 500+ deadlock reports issued
 - ~ 54% of these reports were fixed by developers

Practical Considerations

- The tool optimises for **actionability**, not pure soundness:
 - Prefers fewer, high-quality warnings
 - Accepts some false negatives to keep the noise low
- Some non-fixed reports:
 - May still be real bugs (fixed elsewhere or considered low priority)
 - Some are false positives (e.g. infeasible concurrency patterns)
- Overall: evidence that the analysis finds real, impactful bugs at scale.

Conclusion and Related Work

Related Work

- **Automata-theoretic** approaches
- **Static analyses** for deadlocks
- **Dynamic and hybrid** techniques
- This paper:
 - Places a compositional static analysis with strong theory into this landscape
 - Focuses on large-scale industrial deployment

Takeaways

- A new, **critical-pair based** characterisation of deadlocks for a balanced lock language.
- Deadlock detection in this setting is **decidable** and in **NP**.
- A compositional implementation scaled to tens of millions of LoC in production.
- **Formalisation:** full development mechanised in Coq ($\sim 8.7k$ LOC).

Future Work

Future Directions

- Extend the theory to richer languages (recursion, deterministic guards, nested parallelism)
- Sharpen complexity bounds (e.g. NP-completeness)
- Explore similar compositional ideas for other concurrency bugs

Thank you!

Questions?