

A Compositional Deadlock Detector for Android Java

Darius Mureşan Henrique Luz Rui Xavier

November 15, 2025

Aarhus University
Department of Computer Science

Motivation

Why Deadlocks Matter

- Concurrency is everywhere in Android apps:
 - UI thread + many background threads
 - Callback-driven, event-based execution
 - Shared mutable state (objects used as locks)
- A **deadlock** freezes part (or all) of the app:
 - UI becomes unresponsive
 - Background tasks never finish
 - Users perceive the app as *broken*
- Avoiding and detecting deadlocks is a core reliability problem.

Industrial Setting: Facebook Android Apps

- Target: large Android apps under continuous development:
 - Tens of millions of lines of code (LoC)
 - Thousands of revisions per day
- Goal of the analysis:
 - Run at code-review time on every commit
 - Give feedback to developers **within minutes**
- This rules out whole-program, from-scratch analyses for every change.

Key Challenges

- **Scalability:**
 - Cannot re-analyse the whole app for each commit
 - Need a compositional, incremental analysis
- **Usefulness for developers:**
 - Too many false positives ⇒ warnings are ignored
 - The focus is on **actionable reports**, not proving absence of deadlocks
- **The research question:** can we get both **theoretical guarantees** and **practical performance** at this scale?

Research Gap and Question

Limitations of Existing Approaches

- Many existing deadlock analyses:
 - Assume access to the **whole program**
 - Do not support compositional, change-focused analysis
- Tools often prioritise:
 - Strong soundness guarantees
 - But with high false-positive rates in practice
- For huge codebases, this is misaligned with developer needs:
 - Developers prefer fewer, more trustworthy warnings
 - It is acceptable to miss some rare deadlocks

Research Question and Contributions

Main Question

Can we design a **compositional** static analysis for **Android Java** that detects deadlocks in large codebases, with:

- A clean **theoretical characterisation** of deadlocks
- A **decidable** and **tractable** core problem
- A practical implementation integrated in CI

High-Level Contributions

- Abstract language with balanced re-entrant locks and nondeterministic control
- New deadlock condition based on **critical pairs** of threads
- Proof that deadlock detection in this language is decidable and in NP
- Compositional implementation in Facebook's INFER, deployed

Critical Pairs and Deadlock Detection

Intuition: Critical Pairs

Critical Pair (Informal)

A **critical pair** of a thread is a pair (X, ℓ) such that some execution of the thread acquires an *unheld* lock ℓ while already holding exactly the set of locks X .

- Captures **which locks are held** when a new lock is acquired.
- Abstracts away the concrete control-flow and interleavings.
- Each thread has a **finite** set of critical pairs.

Two-Thread Deadlock Condition

Deadlock Condition (2 Threads)

For two threads C_1 and C_2 :

$$C_1 \parallel C_2 \text{ deadlocks} \iff \exists (X_1, \ell_1) \in \text{Crit}(C_1), (X_2, \ell_2) \in \text{Crit}(C_2)$$

such that

$$\ell_1 \in X_2, \quad \ell_2 \in X_1, \quad X_1 \cap X_2 = \emptyset.$$

- Each thread holds a lock the other is trying to acquire.
- The currently held lock sets do not overlap.
- Generalises to arbitrarily many threads (Theorem 4.4 in the paper).

Example: Classic Two-Thread Deadlock

Two Threads

$C_1 : acq(x); acq(y); skip; rel(y); rel(x)$

$C_2 : acq(y); acq(x); skip; rel(x); rel(y)$

- Critical pairs:

$$\text{Crit}(C_1) = \{(\emptyset, x), (\{x\}, y)\}$$

$$\text{Crit}(C_2) = \{(\emptyset, y), (\{y\}, x)\}$$

- Take $(X_1, \ell_1) = (\{x\}, y)$ and $(X_2, \ell_2) = (\{y\}, x)$:
 - $\ell_1 = y \in X_2 = \{y\}$
 - $\ell_2 = x \in X_1 = \{x\}$
 - $X_1 \cap X_2 = \emptyset$
- Condition holds \Rightarrow deadlock is possible.

Guard Locks: Breaking the Deadlock

Adding a Guard Lock

$$C'_1 = \text{acq}(z); C_1; \text{rel}(z)$$

$$C'_2 = \text{acq}(z); C_2; \text{rel}(z)$$

- Now:

$$\text{Crit}(C'_1) = \{(\emptyset, z), (\{z\}, x), (\{z, x\}, y)\}$$

$$\text{Crit}(C'_2) = \{(\emptyset, z), (\{z\}, y), (\{z, y\}, x)\}$$

- Any potentially conflicting critical pairs share lock z :

$$\{z, x\} \cap \{z, y\} = \{z\} \neq \emptyset$$

- Deadlock condition fails \Rightarrow no deadlock.
- Intuition: z acts as a *guard lock* protecting the region.

Concurrent Programs

Abstract Language for Concurrency

- Abstracts away data and heap; focuses only on locks and control.
- Statements C are built from:

$$\begin{aligned} C ::= & \text{skip} \mid p() \mid \text{acq}(\ell) \mid \text{rel}(\ell) \mid C; C \\ & \mid \text{if}(*) \text{ then } C \text{ else } C \mid \text{while}(*) \text{ do } C \end{aligned}$$

- **Balanced statements** enforce lock discipline:
 - Locks must be released in LIFO order
 - Models scoped constructs like synchronized and `std::lock_guard`
- A parallel program is an n -tuple $C_1 \parallel \dots \parallel C_n$.

Lock States and Configurations

- Locks are **re-entrant**:
 - A thread can acquire the same lock multiple times
 - Lock state $L : \text{Locks} \rightarrow \mathbb{N}$ counts acquisitions
- We write $[L] = \{\ell \mid L(\ell) > 0\}$ for held locks.
- A configuration is a pair $\langle C, L \rangle$:
 - C : current statement
 - L : lock state
- A concurrent configuration:

$$\langle C_1 || \dots || C_n, (L_1, \dots, L_n) \rangle$$

where each thread has its own lock state.

Deadlock Definition

- Sequential step: $\langle C, L \rangle \rightarrow \langle C', L' \rangle$.
- Parallel step (rule (PAR I)): advance one thread, provided it does not grab a lock already held by some other thread.
- Intuitively, we are deadlocked when:
 - Each thread can individually make a step (sequentially)
 - But no *parallel* step is possible any more
- Formal definition (simplified):
 - A concurrent configuration is deadlocked if at least two threads are stuck in this way.
 - A program deadlocks if some reachable configuration is deadlocked.

Program Execution Traces

Executions as Traces

- Each step either:
 - Leaves the lock state unchanged
 - Acquires a lock ℓ
 - Releases a lock ℓ
- We record only lock actions as a **trace** over alphabet Σ :

$$\Sigma = \{\ell \mid \ell \in \text{Locks}\} \cup \{\bar{\ell} \mid \ell \in \text{Locks}\}$$

- Example:
 - Statement: `acq(x); if(*) then acq(y); rel(y) else acq(z); rel(z); rel(x)`
 - Possible traces: `xy\yx` and `xz\zx`

Dyck Words and Balanced Locking

- For balanced statements, traces are **Dyck words**:
 - Well-parenthesised strings of opens and closes
 - Locks always released in reverse order of acquisition
- Key property:
 - For any balanced statement C , all traces in $L(C)$ are Dyck words.
 - Executions never underflow the lock stack.
- This structure is crucial for the decidability and the critical-pair characterisation of deadlocks.

Languages of Statements

- Each statement C defines a language $L(C) \subseteq \Sigma^*$:
 - All traces of possible executions of C
 - Defined inductively:

$$L(\text{skip}) = \{\varepsilon\}$$

$$L(p())$$

$$L(\text{acq}(\ell)) = \{\ell\}$$

$$L(\text{rel}(\ell))$$

$$L(C_1; C_2) = L(C_1) \cdot L(C_2) \quad L(\text{if } (*) \text{ then } C_1 \text{ else } C_2)$$

$$L(\text{while } (*) \text{ do } C) = L(C)^*$$

- For balanced C , $L(C)$ is regular and consists of Dyck words.

Soundness and Completeness

Critical Pairs: Formal Definition

Critical Pairs of a Statement

For a balanced statement C :

$$\text{Crit}(C) = \{(\lfloor \langle u \rangle \rfloor, \ell) \mid \exists v. \, ulv \in L(C) \text{ and } \ell \notin \lfloor \langle u \rangle \rfloor\}$$

where $\langle u \rangle$ is the cumulative lock effect of trace u .

- This is equivalent to: *there exists an execution that acquires ℓ while holding exactly X .*
- The paper shows the execution-based and language-based definitions coincide.

Deadlock Characterisation Theorem

Theorem 4.4 (Simplified)

A parallel program $C_1 \parallel \dots \parallel C_n$ deadlocks iff there exists an index set $I \subseteq \{1, \dots, n\}$ with $|I| \geq 2$ and critical pairs $(X_i, \ell_i) \in \text{Crit}(C_i)$ for each $i \in I$ such that:

$$X_i \cap \bigcup_{j \neq i} X_j = \emptyset \quad \text{and} \quad \ell_i \in \bigcup_{j \neq i} X_j \quad \text{for all } i \in I.$$

- Each thread holds locks needed by the others.
- Held-lock sets are pairwise disjoint.
- This condition is both **sound** and **complete**.

Proof Idea (Very High Level)

- Direction (\Rightarrow) (deadlock \Rightarrow critical-pair conflict):
 - Start from a deadlocked configuration.
 - Project onto each thread and analyse the last lock acquisition step.
 - Show that the corresponding lock states form the required X_i .
- Direction (\Leftarrow) (critical-pair conflict \Rightarrow deadlock):
 - Assume such critical pairs exist.
 - Construct executions for each thread that reach the corresponding states.
 - Use a scheduling argument to show a global deadlocked configuration exists.
- Balanced locking and Dyck-word reasoning are crucial in both directions.

Complexity

Computing Critical Pairs Compositionally

- The paper gives equations (C1)–(C6) describing $\text{Crit}(C)$ by syntax of C .
- Examples:

$$\text{Crit}(\text{skip}) = \emptyset$$

$$\text{Crit}(p()) = \text{Crit}(\text{body}(p))$$

$$\text{Crit}(C; C') = \text{Crit}(C) \cup \text{Crit}(C')$$

$$\text{Crit}(\text{if}(* \text{ then } C \text{ else } C') = \text{Crit}(C) \cup \text{Crit}(C')$$

$$\text{Crit}(\text{acq}(\ell); C; \text{rel}(\ell)) = \{(\emptyset, \ell)\} \cup \{(X \cup \{\ell\}, \ell') \mid (X, \ell') \in$$

- These identities allow a bottom-up computation of $\text{Crit}(C)$.

Complexity Bounds

- **Finite and computable:**
 - For any balanced C , $\text{Crit}(C)$ is finite.
- Complexity of deadlock detection (Theorem 5.5):
 - The deadlock problem for this language is **decidable** and lies in **NP**.
 - Idea: nondeterministically guess a set of threads and critical pairs, then check the deadlock condition in polynomial time.
- For programs without procedure calls:
 - Computing $\text{Crit}(C)$ is polynomial in program size.
- Lower bounds (e.g. NP-completeness) are left as future work.

Implementation

Core Abstract Interpretation

- Implementation is an abstract interpretation that computes critical pairs.
- Abstract state:

$$\alpha = \langle L, Z \rangle$$

where

- L : abstract lock state
- Z : set of (approximate) critical pairs
- Each command C defines a transformer $\llbracket C \rrbracket(\alpha)$.
- The join operation on states:

$$\langle L, Z_1 \rangle \sqcup \langle L, Z_2 \rangle = \langle L, Z_1 \cup Z_2 \rangle$$

Compositionality of the Analysis

- Procedure calls are handled via **summaries**:
 - For each procedure p we precompute $\text{Crit}(\text{body}(p))$.
 - At a call $p()$, we combine the current state with p 's summary.
- Consequences:
 - When code changes, only affected procedures and their callers need re-analysis.
 - Most of the program can be reused from previous runs.
 - This is essential for deployment in continuous integration.

Adapting to Android Java

- The abstract language is mapped to real Java/Android code:
 - synchronized methods/blocks ⇒ balanced lock regions
 - Re-entrant monitors modelled as nested acquisitions
- Android-specific refinements:
 - Partial path-sensitivity for methods like tryLock() and UI-thread checks
 - Lock naming via access-paths (this.f.g, etc.)
 - Thread identity domain (@UiThread, @WorkerThread, background)
- Implemented as the starvation analyser inside INFER.

Deployment and Results

Integration in Facebook's CI

- INFER is part of Facebook's continuous integration:
 - Every Android commit triggers static analyses, including deadlock analysis.
 - The analyser appears as an automated reviewer on code reviews.
- The deadlock analysis targets **code changes**, not whole apps:
 - Summarise modified methods and their dependents
 - Use heuristics to find relevant methods that share locks

Quantitative Impact

- Deployed on all Android commits for ~ 2 years.
- Scale:
 - Hundreds of thousands of commits analysed
 - Typically 2k – 5k methods per commit
- Performance:
 - Median total analysis time ≈ 90 seconds per commit
- Effectiveness:
 - 500+ deadlock reports issued
 - $\sim 54\%$ of these reports were fixed by developers

Practical Considerations

- The tool optimises for **actionability**, not pure soundness:
 - Prefers fewer, high-quality warnings
 - Accepts some false negatives to keep the noise low
- Some non-fixed reports:
 - May still be real bugs (fixed elsewhere or considered low priority)
 - Some are false positives (e.g. infeasible concurrency patterns)
- Overall: evidence that the analysis finds real, impactful bugs at scale.

Conclusion and Related Work

Related Work (Very Brief)

- **Automata-theoretic** approaches:
 - Communicating pushdown systems, nested locks, LTL model checking
- **Static analyses** for deadlocks:
 - Path-insensitive and path-sensitive tools for Java and C
 - Often whole-program and non-compositional
- **Dynamic and hybrid** techniques:
 - Runtime monitoring, lock-order graphs, concolic execution
- This paper:
 - Places a compositional static analysis with strong theory into this landscape
 - Focuses on large-scale industrial deployment

Takeaways and Future Work

- A new, **critical-pair based** characterisation of deadlocks for a balanced lock language.
- Deadlock detection in this setting is **decidable** and in **NP**.
- A compositional implementation scaled to tens of millions of LoC in production.
- **Formalisation:** full development mechanised in Coq ($\sim 8.7k$ LOC).

Future Directions

- Extend the theory to richer languages (recursion, deterministic guards, nested parallelism)
- Sharpen complexity bounds (e.g. NP-completeness)
- Explore similar compositional ideas for other concurrency bugs

Thank you!

Questions?