

# A Compositional Deadlock Detector for Android Java

---

Darius Mureşan    Henrique Luz    Rui Xavier

November 19, 2025

Aarhus University  
Department of Computer Science

# Motivation

---

# Deadlocks: A Refresher

- A **deadlock** occurs when two or more threads are each waiting for the other to release a lock, so none can make progress.

## Example (two threads, two locks)

Thread 1: acq(x); acq(y);

Thread 2: acq(y); acq(x);

- One possible execution:
  1. Thread 1 acquires x
  2. Thread 2 acquires y
  3. Thread 1 now waits for y, Thread 2 waits for x
- Neither thread can continue  $\implies$  **deadlock**.

# Why Deadlocks Matter

- Concurrency is everywhere in Android apps:
  - UI thread + many background threads
  - Callback-driven, event-based execution
  - Shared mutable state (objects used as locks)
- A **deadlock** freezes part (or all) of the app:
  - UI becomes unresponsive
  - Background tasks never finish
  - Users perceive the app as *broken*
- Avoiding and detecting deadlocks is a core reliability problem.

# Industrial Setting: Facebook Android Apps

- Target: large Android apps under continuous development:
  - Tens of millions of lines of code (LoC)
  - Thousands of revisions per day
- Goal of the analysis:
  - Run at code-review time on every commit
  - Give feedback to developers **within minutes**
- This rules out whole-program, from-scratch analyses for every change.

# Key Challenges

---

- **Scalability:**
  - Cannot re-analyse the whole app for each commit
  - Need a compositional, incremental analysis
- **Usefulness for developers:**
  - Too many false positives ⇒ warnings are ignored
  - The focus is on **actionable reports**, not proving absence of deadlocks
- **The research question:** can we get both **theoretical guarantees** and **practical performance** at this scale?

## Research Gap and Question

---

# Limitations of Existing Approaches

---

- Many existing deadlock analyses:
  - Assume access to the **whole program**
  - Do not support compositional, change-focused analysis
- Tools often prioritise:
  - Strong soundness guarantees
  - But with high false-positive rates in practice
- For huge codebases, this is misaligned with developer needs:
  - Developers prefer fewer, more trustworthy warnings
  - It is acceptable to miss some rare deadlocks

# Research Question

---

## Main Question

Can we design a **compositional** static analysis for **Android Java** that detects deadlocks in large codebases, with:

- A clean **theoretical characterisation** of deadlocks
- A **decidable** and **tractable** core problem
- A practical implementation integrated in CI

# High-Level Contributions

---

## High-Level Contributions

- Abstract language with balanced re-entrant locks and nondeterministic control
- New deadlock condition based on **critical pairs** of threads
- Proof that deadlock detection in this language is decidable and in NP
- Compositional implementation in Facebook's INFER, deployed at scale

## Concurrent Programs

---

# Abstract Language for Concurrency

- Abstracts away data and heap; focuses only on locks and control.
- Statements  $C$  are built from:

$$\begin{aligned} C ::= & \text{skip} \mid p() \mid acq(\ell); C; rel(\ell) \mid C; C \\ & \mid \text{if}(*) \text{ then } C \text{ else } C \mid \text{while}(*) \text{ do } C \end{aligned}$$

- **Balanced statements** enforce lock discipline:
  - Locks must be released in LIFO order
  - Models scoped constructs like synchronized and `std::lock_guard`
- A parallel program is an  $n$ -tuple  $C_1 \parallel \dots \parallel C_n$ .

# Example Program in the Abstract Language

## Example

$$C = acq(x); (\text{if } (*) \text{ then } acq(y); rel(y) \text{ else skip}); rel(x)$$

- This program:
  - Always acquires and releases  $x$
  - Nondeterministically either:
    - acquires and releases  $y$ , or
    - does nothing in the branch
- Shows how nested acquisitions map to lock scopes in Java.

# Lock States and Configurations

- Locks are **re-entrant**:
  - A thread can acquire the same lock multiple times
  - Lock state  $L : \text{Locks} \rightarrow \mathbb{N}$  counts acquisitions
- We write  $[L] = \{\ell \mid L(\ell) > 0\}$  for held locks.
- A configuration is a pair  $\langle C, L \rangle$ :
  - $C$ : current statement
  - $L$ : lock state
- A concurrent configuration:

$$\langle C_1 || \dots || C_n, (L_1, \dots, L_n) \rangle$$

where each thread has its own lock state.

## Deadlock Definition (1)

- Sequential step:  $\langle C, L \rangle \rightarrow \langle C', L' \rangle$ .

### Example

$\langle \text{acq}(x); \text{skip}; \text{rel}(x), L \rangle \rightarrow \langle \text{skip}; \text{rel}(x), L[x \text{ } ++] \rangle$

(acquiring lock  $x$  increases its count)

- Parallel step: advance one thread, provided it does not grab a lock already held by another thread.
- Intuitively, we are deadlocked when:
  - Each thread can individually make a step (sequential)
  - But no *parallel* step is possible any more

## Deadlock Definition (2)

---

- Formal definition (simplified):
  - A concurrent configuration is deadlocked if at least two threads are stuck in this way.
  - A program deadlocks if some reachable configuration is deadlocked.

# **Critical Pairs and Deadlock Detection**

---

## Critical Pair (Informal)

A **critical pair** of a thread is a pair  $(X, \ell)$  such that some execution of the thread acquires an *unheld* lock  $\ell$  while already holding exactly the set of locks  $X$ .

- Captures **which locks are held** when a new lock is acquired.
- Abstracts away the concrete control-flow and interleavings.
- Each thread has a **finite** set of critical pairs.

# Critical Pairs: A Simple Example

## Statement

$$C = \text{acq}(x); \text{ acq}(y); \text{ rel}(y); \text{ rel}(x)$$

## Computing $\text{Crit}(C)$

$$\text{Crit}(C) = \{ (\emptyset, x), (\{x\}, y) \}$$

- At the first acquisition of  $x$ : no locks held  $\Rightarrow (\emptyset, x)$
- Later, the thread holds  $\{x\}$  when acquiring  $y \Rightarrow (\{x\}, y)$

## Two-Thread Deadlock Condition

### Deadlock Condition (2 Threads)

For two statements  $C_1$  and  $C_2$  running concurrently:

$$C_1 \parallel C_2 \text{ deadlocks} \iff \exists (X_1, \ell_1) \in \text{Crit}(C_1), (X_2, \ell_2) \in \text{Crit}(C_2)$$

such that

$$\ell_1 \in X_2, \quad \ell_2 \in X_1, \quad X_1 \cap X_2 = \emptyset.$$

- Each thread holds a lock the other is trying to acquire.
- The currently held lock sets do not overlap.
- Generalises to arbitrarily many threads.

## Example: Classic Two-Thread Deadlock

### Two Threads

$C_1 : acq(x); acq(y); skip; rel(y); rel(x)$

$C_2 : acq(y); acq(x); skip; rel(x); rel(y)$

- Critical pairs:

$$\text{Crit}(C_1) = \{(\emptyset, x), (\{x\}, y)\}$$

$$\text{Crit}(C_2) = \{(\emptyset, y), (\{y\}, x)\}$$

- Take  $(X_1, \ell_1) = (\{x\}, y)$  and  $(X_2, \ell_2) = (\{y\}, x)$ :
  - $\ell_1 = y \in X_2 = \{y\}$
  - $\ell_2 = x \in X_1 = \{x\}$
  - $X_1 \cap X_2 = \emptyset$
- Condition holds  $\Rightarrow$  deadlock is possible.

# Guard Locks: Breaking the Deadlock

## Adding a Guard Lock

$$C'_1 = \text{acq}(z); C_1; \text{rel}(z)$$

$$C'_2 = \text{acq}(z); C_2; \text{rel}(z)$$

- Now:

$$\text{Crit}(C'_1) = \{(\emptyset, z), (\{z\}, x), (\{z, x\}, y)\}$$

$$\text{Crit}(C'_2) = \{(\emptyset, z), (\{z\}, y), (\{z, y\}, x)\}$$

- Any potentially conflicting critical pairs share lock  $z$ :

$$\{z, x\} \cap \{z, y\} = \{z\} \neq \emptyset$$

- Deadlock condition fails  $\Rightarrow$  no deadlock.
- Intuition:  $z$  acts as a *guard lock* protecting the region.

# Program Execution Traces

---

# Executions as Traces

- Each step either:
  - Leaves the lock state unchanged
  - Acquires a lock  $\ell$
  - Releases a lock  $\ell$
- We record only lock actions as a **trace** over alphabet  $\Sigma$ :

$$\Sigma = \{\ell \mid \ell \in \text{Locks}\} \cup \{\bar{\ell} \mid \ell \in \text{Locks}\}$$

## Example

- Statement: `acq(x); if(*) then acq(y); rel(y) else acq(z); rel(z); rel(x)`
- Possible traces:  $xy\overline{yx}$  and  $xz\overline{zx}$

## Executions as Traces (Examples)

### Example Statement

$$C = \text{acq}(x); \text{acq}(y); \text{rel}(y); \text{rel}(x)$$

- Valid trace:

$x y \bar{y} \bar{x}$

Let us now consider:

$$C' = \text{acq}(x); \text{acq}(y); \text{rel}(x); \text{rel}(y)$$

- Invalid trace (bad nesting):

$x y \bar{x} \bar{y} \Rightarrow \text{lock } x \text{ released while } y \text{ still held}$

# Dyck Words and Balanced Locking

---

- For balanced statements, traces are **Dyck words**:
  - Well-parenthesised strings of opens and closes
  - Locks always released in reverse order of acquisition
- Key property:
  - For any balanced statement  $C$ , all traces in  $L(C)$  are Dyck words.
  - Executions never underflow the lock stack.
- This structure is crucial for the decidability and the critical-pair characterisation of deadlocks.

## Dyck Words and Balanced Locking (Example)

- For balanced statements, traces form **Dyck words**:
  - Perfectly nested lock acquisitions / releases
  - The lock stack never underflows

### Valid Dyck Word

$x \, y \, \bar{y} \, z \, \bar{z} \, \bar{x}$

- Corresponds to nested Acquire–Release behaviour:

$acq(x); \, acq(y); \, rel(y); \, acq(z); \, rel(z); \, rel(x)$

### Invalid Word

$x \, y \, \bar{x} \, \bar{y}$

- Releases  $x$  before  $y$ : violates LIFO discipline.

# Languages of Statements

- Each statement  $C$  defines a language  $L(C) \subseteq \Sigma^*$ :
  - All traces of possible executions of  $C$
  - Defined inductively:

$$L(\text{skip}) = \{\varepsilon\}$$

$$L(\text{rel}(\ell)) = \{\bar{\ell}\}$$

$$L(p()) = L(\text{body}(p))$$

$$L(C_1; C_2) = L(C_1) \cdot L(C_2)$$

$$L(\text{acq}(\ell)) = \{\ell\}$$

$$L(\text{while}(\ast) \text{ do } C) = L(C)^*$$

$$L(\text{if}(\ast) \text{ then } C_1 \text{ else } C_2) = L(C_1) \cup L(C_2)$$

- For balanced  $C$ ,  $L(C)$  is regular and consists of Dyck words.

## Soundness and Completeness

---

## Cumulative Lock Effect of a Trace

---

- For a trace  $u \in \Sigma^*$ ,  $\langle u \rangle$  denotes its **cumulative lock effect**:
  - Start from the empty lock multiset
  - Read actions in  $u$  left to right
  - Each  $x$  increments the count of lock  $x$
  - Each  $\bar{x}$  decrements the count of lock  $x$
- The set of *currently held locks* after reading  $u$  is:

$$[\langle u \rangle] = \{ \ell \mid \langle u \rangle(\ell) > 0 \}.$$

## Example

Trace prefix:  $u = x y \bar{y}$

$$\langle \varepsilon \rangle = \emptyset$$

$$\langle x \rangle = \{x\}$$

$$\langle xy \rangle = \{x, y\}$$

$$\langle xy\bar{y} \rangle = \{x\}$$

So:

$$\lfloor \langle u \rangle \rfloor = \{x\}.$$

## Critical Pairs: Formal Definition

### Critical Pairs of a Statement

For a balanced statement  $C$ :

$$\text{Crit}(C) = \{(\lfloor \langle u \rangle \rfloor, \ell) \mid \exists v. \, ulv \in L(C) \text{ and } \ell \notin \lfloor \langle u \rangle \rfloor\}$$

where  $\langle u \rangle$  is the cumulative lock effect of trace  $u$ .

- This is equivalent to: *there exists an execution that acquires  $\ell$  while holding exactly  $X$ .*
- The paper shows the execution-based and language-based definitions coincide.

## Critical Pairs: Formal Definition (Example)

### Example Trace

$$u\ell v = x \ y \ z \ \bar{z} \ \bar{y} \ \bar{x}$$

- Take  $u = x \ y$ ,  $\ell = z$ :

$$\langle u \rangle = \{x, y\}$$

- Since  $z \notin \{x, y\}$ , acquiring  $z$  yields:

$$(\{x, y\}, z) \in \text{Crit}(C)$$

- Meaning: During some execution, the thread acquires  $z$  while holding exactly  $\{x, y\}$ .

# Deadlock Characterisation Theorem

## Theorem 4.4 (Simplified)

A parallel program  $C_1||\dots||C_n$  deadlocks iff there exists an index set  $I \subseteq \{1, \dots, n\}$  with  $|I| \geq 2$  and critical pairs  $(X_i, \ell_i) \in \text{Crit}(C_i)$  for each  $i \in I$  such that:

$$X_i \cap \bigcup_{j \neq i} X_j = \emptyset \quad \text{and} \quad \ell_i \in \bigcup_{j \neq i} X_j \quad \text{for all } i \in I.$$

- Each thread holds locks needed by the others.
- Held-lock sets are pairwise disjoint.
- This condition is both **sound** and **complete**.

# Complexity

---

## Computing Critical Pairs Compositionally

- The paper gives 6 equations describing  $\text{Crit}(C)$  by syntax of  $C$ .
- Examples:

$$\text{Crit}(\text{skip}) = \emptyset$$

$$\text{Crit}(p()) = \text{Crit}(\text{body}(p))$$

$$\text{Crit}(C; C') = \text{Crit}(C) \cup \text{Crit}(C')$$

$$\text{Crit}(\text{if } (*) \text{ then } C \text{ else } C') = \text{Crit}(C) \cup \text{Crit}(C')$$

$$\text{Crit}(\text{acq}(\ell); C; \text{rel}(\ell)) =$$

$$\{(\emptyset, \ell)\} \cup \{(X \cup \{\ell\}, \ell') \mid (X, \ell') \in \text{Crit}(C), \ell' \neq \ell\}$$

- These identities allow a bottom-up computation of  $\text{Crit}(C)$ .

# Complexity Bounds

---

- **Finite and computable:**
  - For any balanced  $C$ ,  $\text{Crit}(C)$  is finite.
- Complexity of deadlock detection:
  - The deadlock problem for this language is **decidable** and lies in **NP**.
  - Idea: nondeterministically guess a set of threads and critical pairs, then check the deadlock condition in polynomial time.
- For programs without procedure calls:
  - Computing  $\text{Crit}(C)$  is polynomial in program size.
- Lower bounds (e.g. NP-completeness) are left as future work.

## **Implementation**

---

# Core Abstract Interpretation

- Implementation is an abstract interpretation that computes critical pairs.
- Abstract state:

$$\alpha = \langle L, Z \rangle$$

where

- $L$ : abstract lock state
- $Z$ : set of (approximate) critical pairs
- Each command  $C$  defines a transformer  $\llbracket C \rrbracket(\alpha)$ .
- The join operation on states:

$$\langle L, Z_1 \rangle \sqcup \langle L, Z_2 \rangle = \langle L, Z_1 \cup Z_2 \rangle$$

# Compositionality of the Analysis

---

- Procedure calls are handled via **summaries**:
  - For each procedure  $p$  we precompute  $\text{Crit}(\text{body}(p))$ .
  - At a call  $p()$ , we combine the current state with  $p$ 's summary.
- Consequences:
  - When code changes, only affected procedures and their callers need re-analysis.
  - Most of the program can be reused from previous runs.
  - This is essential for deployment in continuous integration.

# Adapting to Android Java

---

- The abstract language is mapped to real Java/Android code:
  - synchronized methods/blocks ⇒ balanced lock regions
  - Re-entrant monitors modelled as nested acquisitions
- Android-specific refinements:
  - Partial path-sensitivity for methods like tryLock() and UI-thread checks
  - Lock naming via access-paths (this.f.g, etc.)
  - Thread identity domain (@UiThread, @WorkerThread, background)
- Implemented as the starvation analyser inside INFER.

## **Deployment and Results**

---

# Integration in Facebook's CI

---

- INFER is part of Facebook's continuous integration:
  - Every Android commit triggers static analyses, including deadlock analysis.
  - The analyser appears as an automated reviewer on code reviews.
- The deadlock analysis targets **code changes**, not whole apps:
  - Summarise modified methods and their dependents
  - Use heuristics to find relevant methods that share locks

## Quantitative Impact

---

- Deployed on all Android commits for ~2 years.
- Scale:
  - Hundreds of thousands of commits analysed
  - Typically 2k – 5k methods per commit
- Performance:
  - Median analysis time of 90 seconds per commit
  - Average analysis time of 213 seconds per commit
- Effectiveness:
  - 500+ deadlock reports issued
  - ~ 54% of these reports were fixed by developers

# Practical Considerations

---

- The tool optimises for **actionability**, not pure soundness:
  - Prefers fewer, high-quality warnings
  - Accepts some false negatives to keep the noise low
- Some non-fixed reports:
  - May still be real bugs (fixed elsewhere or considered low priority)
  - Some are false positives (e.g. infeasible concurrency patterns)
- Overall: evidence that the analysis finds real, impactful bugs at scale.

## **Conclusion and Related Work**

---

## Related Work

---

- **Automata-theoretic** approaches
- **Static analyses** for deadlocks
- **Dynamic and hybrid** techniques
- This paper:
  - Places a compositional static analysis with strong theory into this landscape
  - Focuses on large-scale industrial deployment

## Takeaways

---

- A new, **critical-pair based** characterisation of deadlocks for a balanced lock language.
- Deadlock detection in this setting is **decidable** and in **NP**.
- A compositional implementation scaled to tens of millions of LoC in production.
- **Formalisation:** full development mechanised in Coq ( $\sim 8.7k$  LOC).

# Future Work

---

## Future Directions

- Extend the theory to richer languages (recursion, deterministic guards, nested parallelism)
- Sharpen complexity bounds (e.g. NP-completeness)
- Explore similar compositional ideas for other concurrency bugs

## Pros and Cons

---

## Pros

---

- Achieves a balance between solid theory and real industrial deployment.
- Compositionality is not only formal but practical in CI workflows.
- Critical pairs provide compact, human-interpretable summaries of thread behaviour.
- Abstraction is stable across code evolution, making the analysis maintainable long-term.

## Cons

---

- Compositionality may miss multi-layer deadlocks spread across the call graph (false negatives).
- Strong reliance on balanced locking assumptions; custom sync patterns break the model.
- The first full-program analysis remains expensive despite good per-commit performance.
- The reported statistics aggregate all categories of deadlock reports; separating them would clarify how many were real bugs, low-priority issues, or false positives.
- The evaluation reports total CI analysis time, but does not isolate the cost of *this* analyser, making its standalone overhead harder to assess.

Thank you!

Questions?