

# A Compositional Deadlock Detector for Android Java

---

Darius Mureşan    Henrique Luz    Rui Xavier

November 12, 2025

Aarhus University  
Department of Computer Science

# **Overview**

---

## Overview

---

- We present a **compositional static analysis** for detecting deadlocks in Android Java.
- Designed for **industrial-scale** codebases (tens of millions of lines).
- Implemented in **INFER**, deployed at Facebook for 2+ years.
- Achieves a **54% developer fix rate**.

# Motivation

---

## Why Study Deadlocks?

---

- Deadlocks are a key challenge in concurrent programming.
- Occur when threads **cyclically wait on each other's locks**.
- Result: entire system halts — loss of responsiveness, reliability.
- Especially relevant for **Android apps** using Java's **synchronized blocks**.

## Classical Example

---

### Dijkstra's Dining Philosophers

Five philosophers share forks and a bowl of spaghetti. Each must hold two forks to eat — potential for a circular wait.

- Illustrates the essential nature of deadlock: mutual waiting.
- Analogous to threads waiting for locks.

# Deadlocks in Industry

---

- At Facebook, Android apps exceed **10M+ lines of code**.
- Thousands of commits daily ⇒ rapid iteration.
- Developers need feedback in **under 15 minutes**.
- Whole-program analyses are too slow and memory-intensive.

## Goal

**Fast, scalable, accurate deadlock detection** integrated into CI pipelines.

## Challenges in Traditional Analyses

---

- Reanalyze entire program on each commit.
- Poor scalability and high false-positive rates.
- Lack **compositionality**.
- Need: an analysis that focuses only on changed code and its dependencies.

## **Research Gap and Question**

---

# The Research Gap

---

- Existing tools are **non-compositional**.
- They often trade soundness for scalability.
- We need a **mathematically grounded, incremental** analysis.

## Core Challenge

Can we design a deadlock detector that is both **sound in theory** and **scalable in practice**?

# Research Question

---

## Main Question

Can we develop a **compositional deadlock detector** for **Android Java** that is sound, complete, and efficient for large industrial codebases?

- Model Java concurrency with **balanced, re-entrant locks**.
- Detect deadlocks via **critical pairs**.
- Integrate into **INFER** for continuous integration.

# **Critical Pairs and Deadlock Detection**

---

# Critical Pairs

Rui ► Is this too much text for one slide? ◀

## Definition

A **critical pair** of a thread is a pair  $(X, \ell)$  such that some execution of the thread acquires an unheld lock  $\ell$  while already holding the set of locks  $X$ .

## Deadlock Condition

For the case of two threads:

$C1 \parallel C2$  deadlocks  $\iff$  there are critical pairs  $(X_1, \ell_1)$  and  $(X_2, \ell_2)$  of  $C1$  and  $C2$  respectively, such that:

$$\ell_1 \in X_2 \quad \text{and} \quad \ell_2 \in X_1 \quad \text{and} \quad X_1 \cap X_2 = \emptyset$$

## Example of Critical Pairs

### Example

Consider the two threads:

$$C1 : acq(x); acq(y); skip; rel(y); rel(x)$$
$$C2 : acq(y); acq(x); skip; rel(x); rel(y)$$

Their critical pairs are:

$$Crit(C1) = \{(\emptyset, x), (\{x\}, y)\}$$

$$Crit(C2) = \{(\emptyset, y), (\{y\}, x)\}$$

Since  $x \in \{x\}$ ,  $y \in \{y\}$  and  $\{x\} \cap \{y\} = \emptyset$ , the deadlock condition holds — **deadlock!**

## Example of Critical Pairs

### Solution

Consider the two threads:

$$C1' : \text{acq}(z); C1; \text{rel}(z)$$

$$C2' : \text{acq}(z); C2; \text{rel}(z)$$

Their critical pairs are:

$$\text{Crit}(C1') = \{(\emptyset, z), (\{z\}, x)\}, (\{z, x\}, y)\}$$

$$\text{Crit}(C2') = \{(\emptyset, z), (\{z\}, y)\}, (\{z, y\}, x)\}$$

Since  $\{z, x\} \cap \{z, y\} = z$ , the deadlock condition does not hold —  
**no deadlock!**

$z$  acts as a guard lock, preventing  $x$  and  $y$  from being accessed by  $C1'$  and  $C2'$  simultaneously.

## Concurrent Programs

---

# Abstract Language

- Simplified model of Java concurrency:

$$\begin{aligned} C ::= & \text{skip} \mid p() \mid \text{acq}(\ell); C; \text{rel}(\ell) \mid C; C \\ & \mid \text{if}(*) \text{ then } C \text{ else } C \mid \text{while}(*) \text{ do } C \end{aligned}$$

- This grammar guarantees balanced pairs of lock acquisitions/releases.
- We consider only non-recursive procedures.
- If  $C$  is balanced, then so is  $\text{body}(p)$  with  $p \in \text{callees}(C)$ .
- A parallel program is an  $n$ -tuple of balanced statements  $C_1 \parallel C_2 \parallel \dots \parallel C_n$ .

# Program Execution Traces

---

## Lock States

---

- We treat locks as **re-entrant and scoped**.
- A *lock state* is a function that maps each lock to its acquisition count.
- A *configuration* is a pair  $\langle C, L \rangle$ , where  $C$  is a statement and  $L$  is the lock state.
- A *concurrent configuration* is then, a pair of the form  $\langle C_1 || \dots || C_n, (L_1, \dots, L_n) \rangle$ .

## Executions as Traces

---

- Executions can be represented as strings of lock acquisitions/releases — e.g. ‘x y y x’.
- Balanced programs produce **Dyck words** (well-nested parentheses).
- Captures the essential locking behaviour.

$$L(C) = \{\text{all possible lock traces of } C\}$$

## Example Trace

### Example

```
acq(x); if(*) then acq(y); rel(y); else acq(z);  
rel(z); rel(x)
```

$$L(C) = \{ xyyx, xzzx \}$$

- Balanced structure guarantees decidability.
- Enables abstract reasoning about possible interleavings.

## Traces as Finite Automata

---

- Each balanced statement can be viewed as a **finite automaton** over lock actions.
- Allows algorithmic computation of lock dependencies.
- Provides the foundation for critical-pair analysis.

## Soundness and Completeness

---

# Deadlock Characterisation

---

## Critical Pair

$(X, \ell)$ : some execution acquires lock  $\ell$  while holding all locks in  $X$ .

- Captures possible lock dependencies.
- Computed for each sequential thread.

## Key Theorem (Simplified)

---

### Theorem 4.4 — Deadlock Condition

Program  $C_1||\dots||C_n$  deadlocks iff there exist critical pairs  $(X_i, \ell_i)$   
s.t.

$$\ell_i \in \bigcup_{j \neq i} X_j \quad \text{and} \quad X_i \cap \bigcup_{j \neq i} X_j = \emptyset$$

**Intuition:** Each thread holds a lock another needs.

# Illustrative Example

---

## Two-Thread Example

C1: acq(x); acq(y); rel(y); rel(x) C2: acq(y);  
acq(x); rel(x); rel(y)

$$\text{Crit}(\text{C1}) = \{(\emptyset, x), (\{x\}, y)\}$$

$$\text{Crit}(\text{C2}) = \{(\emptyset, y), (\{y\}, x)\}$$

Since  $x \in \{y\}$  and  $y \in \{x\}$ , the condition holds — **deadlock!**

# Proof Structure

---

- Define executions → and parallel composition.
- Show equivalence between execution semantics and trace semantics.
- Prove critical-pair condition is **sound** (no missed deadlocks) and **complete** (no spurious ones).

## Result

Existence of deadlock  $\Leftrightarrow$  conflict between critical pairs.

# Complexity

---

## Computing Critical Pairs

---

- Recursive equations (C1–C6) compute  $\text{Crit}(C)$  compositionally.
- Each construct (if, while, seq) has a local combination rule.
- **Example:**

$$\text{Crit}(\text{acq}(\ell); C; \text{rel}(\ell)) = \{(\emptyset, \ell)\} \cup \{(X \cup \{\ell\}, \ell') \mid (X, \ell') \in \text{Crit}(C)\}$$

## Complexity Results

---

- **Finite and computable:**  $\text{Crit}(C)$  always finite.
- Deadlock detection problem is **decidable** and in NP.
- Non-recursive programs  $\Rightarrow$  quadratic time.
- With procedures  $\Rightarrow$  quasi-exponential.

## **Implementation**

---

# Implementation Overview

---

- Implemented as an **abstract interpretation** within INFER.
- Computes method summaries: critical pairs + thread identity.
- Compositionally reuses summaries of unchanged methods.

## Core Idea

Analyse only modified methods and their dependents.

# Abstract State Representation

---

$\alpha = \langle L, Z \rangle$  where  $L$  = lock state,  $Z$  = set of critical pairs

- Join operation:  $\langle L, Z_1 \rangle \sqcup \langle L, Z_2 \rangle = \langle L, Z_1 \cup Z_2 \rangle$
- Each command updates this abstract state.

# Compositionality

---

- Procedure call depends only on:
  1. Current abstract state.
  2. Precomputed summary of the callee.
- Enables incremental reanalysis — ideal for CI/CD.

$$J_p()K\langle L, Z \rangle = \langle L, Z \cup f(L, \text{Crit}(\text{body}(p))) \rangle$$

- **Balanced locking:** uses synchronized.
- **Partial path sensitivity:** e.g., for tryLock() and UI threads.
- **Lock naming:** access-path abstraction (this.f.g, etc.).
- **Thread inference:** uses annotations like @UiThread.

## **Deployment and Results**

---

# Industrial Deployment

---

- Deployed as part of Facebook's **continuous integration system**.
- Runs automatically on every Android commit.
- Appears as an automated "reviewer" commenting on potential deadlocks.

## Results and Impact

---

- Deployed for **2+ years** on all Android commits.
- **500+** deadlock reports issued.
- **54%** of reports fixed by developers.
- Median runtime (all analyses): **90 s per commit.**
- Analyses **2k–5k methods per commit.**

## **Conclusion and Related Work**

---

## Related Work

---

- Builds on automata-theoretic analyses of **pushdown systems**.
- Compared to prior tools:
  - Compositional, not whole-program.
  - Targets balanced re-entrant locks.
  - Prioritizes **actionable results** over completeness.

# Conclusion

- Developed a **sound and complete** compositional analysis.
- Scales to tens of millions of lines.
- Successfully deployed in industry with tangible impact.
- Formalized and proven in **Coq (8.7k LOC)**.

## Future Work

Extend to recursive calls, deterministic control, and nested parallelism.

Thank you!

Questions?